2.Quora_Preprocessing

April 18, 2019

0.0.1 1.2.1 : EDA: Advanced Feature Extraction.

In [10]:
```
import warnings
warnings.filterwarnings("ignore")
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from subprocess import check_output
%matplotlib inline
import plotly.offline as py
py.init_notebook_mode(connected=True)
import plotly.graph_objs as go
import plotly.tools as tls
import os
import gc

import re
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer
from bs4 import BeautifulSoup
from sklearn.preprocessing import MinMaxScaler
from nltk.corpus import stopwords
# This package is used for finding longest common subsequence between two strings
# you can write your own dp code for this
import distance
from nltk.stem import PorterStemmer
from bs4 import BeautifulSoup
from fuzzywuzzy import fuzz
from sklearn.manifold import TSNE
# Import the Required lib packages for WORD-Cloud generation
# https://stackoverflow.com/questions/45625434/how-to-install-wordcloud-in-python3-6
from wordcloud import WordCloud, STOPWORDS
from os import path
from PIL import Image
```

In [11]:
```
#https://stackoverflow.com/questions/12468179/unicodedecodeerror-utf8-codec-cant-deco
if os.path.isfile('df_fe_without_preprocessing_train.csv'):
```

```
        df = pd.read_csv("df_fe_without_preprocessing_train.csv",encoding='latin-1')
        df = df.fillna('')
        df.head()
    else:
        print("get df_fe_without_preprocessing_train.csv from drive or run the previous no
```

In [12]:
```
df.head(2)
```

Out[12]:

	id	qid1	qid2	question1	question2	is_duplicate	freq_qid1	freq_qid2	q1len	q2len	q1_n_words	q2_n_words	word_Common	word_Total	word_share	freq_q1+q2	freq_q1-q2
0	0	1	2	What is the step by step guide to invest in sh...	What is the step by step guide to invest in sh...	0	1	1	66	57	14	12	10.0	23.0	0.434783	2	0
1	1	3	4	What is the story of Kohinoor (Koh-i-Noor) Dia...	What would happen if the Indian government sto...	0	4	1	51	88	8	13	4.0	20.0	0.200000	5	3

3.4 Preprocessing of Text

- Preprocessing:
  - Removing html tags
  - Removing Punctuations
  - Performing stemming
  - Removing Stopwords
  - Expanding contractions etc.

In [13]:
```
# To get the results in 4 decemal points
SAFE_DIV = 0.0001

STOP_WORDS = stopwords.words("english")


def preprocess(x):
    x = str(x).lower()
    x = x.replace(",000,000", "m").replace(",000", "k").replace("", "'").replace("",
                           .replace("won't", "will not").replace("cannot", "can not")
                           .replace("n't", " not").replace("what's", "what is").repla
                           .replace("'ve", " have").replace("i'm", "i am").replace("'
                           .replace("he's", "he is").replace("she's", "she is").repla
                           .replace("%", " percent ").replace("", " rupee ").replace(
                           .replace("", " euro ").replace("'ll", " will")
```

```
    x = re.sub(r"([0-9]+)000000", r"\1m", x)
    x = re.sub(r"([0-9]+)000", r"\1k", x)


    porter = PorterStemmer()
    pattern = re.compile('\W')

    if type(x) == type(''):
        x = re.sub(pattern, ' ', x)


    if type(x) == type(''):
        x = porter.stem(x)
        example1 = BeautifulSoup(x)
        x = example1.get_text()


    return x
```

- Function to Compute and get the features : With 2 parameters of Question 1 and Question 2

3.5 Advanced Feature Extraction (NLP and Fuzzy Features)

Definition: - **Token**: You get a token by splitting sentence a space - **Stop_Word** : stop words as per NLTK. - **Word** : A token that is not a stop_word

Features: - **cwc_min** : Ratio of common_word_count to min lenghth of word count of Q1 and Q2 cwc_min = common_word_count / (min(len(q1_words), len(q2_words)) - **cwc_max** : Ratio of common_word_count to max lenghth of word count of Q1 and Q2 cwc_max = common_word_count / (max(len(q1_words), len(q2_words)) - **csc_min** : Ratio of common_stop_count to min lenghth of stop count of Q1 and Q2 csc_min = common_stop_count / (min(len(q1_stops), len(q2_stops)) - **csc_max** : Ratio of common_stop_count to max lenghth of stop count of Q1 and Q2csc_max = common_stop_count / (max(len(q1_stops), len(q2_stops)) - **ctc_min** : Ratio of common_token_count to min lenghth of token count of Q1 and Q2ctc_min = common_token_count / (min(len(q1_tokens), len(q2_tokens))

- **ctc_max** : Ratio of common_token_count to max lenghth of token count of Q1 and Q2ctc_max = common_token_count / (max(len(q1_tokens), len(q2_tokens))
- **last_word_eq** : Check if First word of both questions is equal or notlast_word_eq = int(q1_tokens[-1] == q2_tokens[-1])
- **first_word_eq** : Check if First word of both questions is equal or notfirst_word_eq = int(q1_tokens[0] == q2_tokens[0])
- **abs_len_diff** : Abs. length differenceabs_len_diff = abs(len(q1_tokens) - len(q2_tokens))
- **mean_len** : Average Token Length of both Questionsmean_len = (len(q1_tokens) + len(q2_tokens))/2
- **fuzz_ratio** : https://github.com/seatgeek/fuzzywuzzy#usage http://chairnerd.seatgeek.com/fuzzywuzzy-fuzzy-string-matching-in-python/

- **fuzz_partial_ratio** : https://github.com/seatgeek/fuzzywuzzy#usage http://chairnerd.seatgeek.com/fuzzywuzzy-fuzzy-string-matching-in-python/
- **token_sort_ratio** : https://github.com/seatgeek/fuzzywuzzy#usage http://chairnerd.seatgeek.com/fuzzywuzzy-fuzzy-string-matching-in-python/
- **token_set_ratio** : https://github.com/seatgeek/fuzzywuzzy#usage http://chairnerd.seatgeek.com/fuzzywuzzy-fuzzy-string-matching-in-python/
- **longest_substr_ratio** : Ratio of length longest common substring to min lenghth of token count of Q1 and Q2longest_substr_ratio = len(longest common substring) / (min(len(q1_tokens), len(q2_tokens))

In [14]:
```
def get_token_features(q1, q2):
    token_features = [0.0]*10

    # Converting the Sentence into Tokens:
    q1_tokens = q1.split()
    q2_tokens = q2.split()

    if len(q1_tokens) == 0 or len(q2_tokens) == 0:
        return token_features
    # Get the non-stopwords in Questions
    q1_words = set([word for word in q1_tokens if word not in STOP_WORDS])
    q2_words = set([word for word in q2_tokens if word not in STOP_WORDS])

    #Get the stopwords in Questions
    q1_stops = set([word for word in q1_tokens if word in STOP_WORDS])
    q2_stops = set([word for word in q2_tokens if word in STOP_WORDS])

    # Get the common non-stopwords from Question pair
    common_word_count = len(q1_words.intersection(q2_words))

    # Get the common stopwords from Question pair
    common_stop_count = len(q1_stops.intersection(q2_stops))

    # Get the common Tokens from Question pair
    common_token_count = len(set(q1_tokens).intersection(set(q2_tokens)))


    token_features[0] = common_word_count / (min(len(q1_words), len(q2_words)) + SAFE
    token_features[1] = common_word_count / (max(len(q1_words), len(q2_words)) + SAFE
    token_features[2] = common_stop_count / (min(len(q1_stops), len(q2_stops)) + SAFE
    token_features[3] = common_stop_count / (max(len(q1_stops), len(q2_stops)) + SAFE
    token_features[4] = common_token_count / (min(len(q1_tokens), len(q2_tokens)) + SA
    token_features[5] = common_token_count / (max(len(q1_tokens), len(q2_tokens)) + SA

    # Last word of both question is same or not
    token_features[6] = int(q1_tokens[-1] == q2_tokens[-1])
```

```
    # First word of both question is same or not
    token_features[7] = int(q1_tokens[0] == q2_tokens[0])

    token_features[8] = abs(len(q1_tokens) - len(q2_tokens))

    #Average Token Length of both Questions
    token_features[9] = (len(q1_tokens) + len(q2_tokens))/2
    return token_features

# get the Longest Common sub string

def get_longest_substr_ratio(a, b):
    if distance.lcsubstrings(a, b) == 0:
        return 0
    else:
        A=distance.lcsubstrings(a, b)
        B=max(min(len(a), len(b)),1)

        return A / B

def extract_features(df):
    # preprocessing each question
    df["question1"] = df["question1"].fillna("").apply(preprocess)
    df["question2"] = df["question2"].fillna("").apply(preprocess)

    print("token features...")

    # Merging Features with dataset

    token_features = df.apply(lambda x: get_token_features(x["question1"], x["questio

    df["cwc_min"]       = list(map(lambda x: x[0], token_features))
    df["cwc_max"]       = list(map(lambda x: x[1], token_features))
    df["csc_min"]       = list(map(lambda x: x[2], token_features))
    df["csc_max"]       = list(map(lambda x: x[3], token_features))
    df["ctc_min"]       = list(map(lambda x: x[4], token_features))
    df["ctc_max"]       = list(map(lambda x: x[5], token_features))
    df["last_word_eq"]  = list(map(lambda x: x[6], token_features))
    df["first_word_eq"] = list(map(lambda x: x[7], token_features))
    df["abs_len_diff"]  = list(map(lambda x: x[8], token_features))
    df["mean_len"]      = list(map(lambda x: x[9], token_features))

    #Computing Fuzzy Features and Merging with Dataset

    # do read this blog: http://chairnerd.seatgeek.com/fuzzywuzzy-fuzzy-string-matchi
    # https://stackoverflow.com/questions/31806695/when-to-use-which-fuzz-function-to
    # https://github.com/seatgeek/fuzzywuzzy
```

```
    print("fuzzy features..")

    df["token_set_ratio"]       = df.apply(lambda x: fuzz.token_set_ratio(x["question
    # The token sort approach involves tokenizing the string in question, sorting the
    # then joining them back into a string We then compare the transformed strings wi
    df["token_sort_ratio"]      = df.apply(lambda x: fuzz.token_sort_ratio(x["questio
    df["fuzz_ratio"]            = df.apply(lambda x: fuzz.QRatio(x["question1"], x["q
    df["fuzz_partial_ratio"]    = df.apply(lambda x: fuzz.partial_ratio(x["question1"
    df["longest_substr_ratio"]  = df.apply(lambda x: get_longest_substr_ratio(x["ques
    return df
```

```
In [15]: if os.path.isfile('nlp_features_train.csv'):
             df = pd.read_csv("nlp_features_train.csv",encoding='latin-1')
             df.fillna('')
         else:
             print("Extracting features for train:")
             df = pd.read_csv("train.csv")
             df = extract_features(df)
             df.to_csv("nlp_features_train.csv", index=False)
         df.head(2)
```

```
Extracting features for train:
token features...
fuzzy features..
```

Out[15]:

	id	qid1	qid2	question1	question2	is_duplicate	cwc_min	cwc_max	csc_min	csc_max	...	ctc_max	last_word_eq	first_word_eq	abs_len_diff	mean_len	token_set_ratio	token_sort_ratio	fuzz_ratio	fuzz_partial_ratio	longest_substr_ratio
0	0	1	2	what is the step by step guide to invest in sh...	what is the step by step guide to invest in sh...	0	0.999980	0.833319	0.999983	0.999983	...	0.785709	0.0	1.0	2.0	13.0	100	93	93	100	1.000000
1	1	3	4	what is the story of kohinoor koh i noor dia...	what would happen if the indian government sto...	0	0.799984	0.399996	0.749981	0.599988	...	0.466664	0.0	1.0	5.0	12.5	86	63	66	75	0.607843

[2 rows x 21 columns]

3.5.1 Analysis of extracted features

3.5.1.1 Plotting Word clouds

- Creating Word Cloud of Duplicates and Non-Duplicates Question pairs
- We can observe the most frequent occuring words

```
In [21]: df_duplicate = df[df['is_duplicate'] == 1]
         dfp_nonduplicate = df[df['is_duplicate'] == 0]

         # Converting 2d array of q1 and q2 and flatten the array: like {{1,2},{3,4}} to {1,2,
         p = np.dstack([df_duplicate["question1"], df_duplicate["question2"]]).flatten()
         n = np.dstack([dfp_nonduplicate["question1"], dfp_nonduplicate["question2"]]).flatten

         print ("Number of data points in class 1 (duplicate pairs) :",len(p))
         print ("Number of data points in class 0 (non duplicate pairs) :",len(n))

         #Saving the np array into a text file
         np.savetxt('train_p.txt', p, delimiter=' ', fmt='%s', encoding ='utf-8')
         np.savetxt('train_n.txt', n, delimiter=' ', fmt='%s', encoding ='utf-8')
```

```
Number of data points in class 1 (duplicate pairs) : 298526
Number of data points in class 0 (non duplicate pairs) : 510054
```

```
In [23]: # reading the text files and removing the Stop Words:
         d = path.dirname('.')

         with open(path.join(d, 'train_p.txt'),'r',encoding='utf-8') as file:
             textp_w = file.read()
         # textp_w = open(path.join(d, 'train_p.txt')).read()
         with open(path.join(d, 'train_n.txt'),'r',encoding='utf-8') as file:
             textn_w = file.read()
         # textn_w = open(path.join(d, 'train_n.txt')).read()
         stopwords = set(STOPWORDS)
         stopwords.add("said")
         stopwords.add("br")
         stopwords.add(" ")
         stopwords.remove("not")

         stopwords.remove("no")
         #stopwords.remove("good")
         #stopwords.remove("love")
         stopwords.remove("like")
         #stopwords.remove("best")
         #stopwords.remove("!")
         print ("Total number of words in duplicate pair questions :",len(textp_w))
         print ("Total number of words in non duplicate pair questions :",len(textn_w))
```

```
Total number of words in duplicate pair questions : 16109886
Total number of words in non duplicate pair questions : 33193067
```

__ Word Clouds generated from duplicate pair question's text __

```
In [24]: wc = WordCloud(background_color="white", max_words=len(textp_w), stopwords=stopwords)
         wc.generate(textp_w)
         print ("Word Cloud for Duplicate Question pairs")
         plt.imshow(wc, interpolation='bilinear')
         plt.axis("off")
         plt.show()
```

```
Word Cloud for Duplicate Question pairs
```

__ Word Clouds generated from non duplicate pair question's text __

```
In [26]: wc = WordCloud(background_color="white", max_words=len(textn_w),stopwords=stopwords)
         # generate word cloud
         wc.generate(textn_w)
         print ("Word Cloud for non-Duplicate Question pairs:")
         plt.imshow(wc, interpolation='bilinear')
         plt.axis("off")
         plt.show()
```

```
Word Cloud for non-Duplicate Question pairs:
```

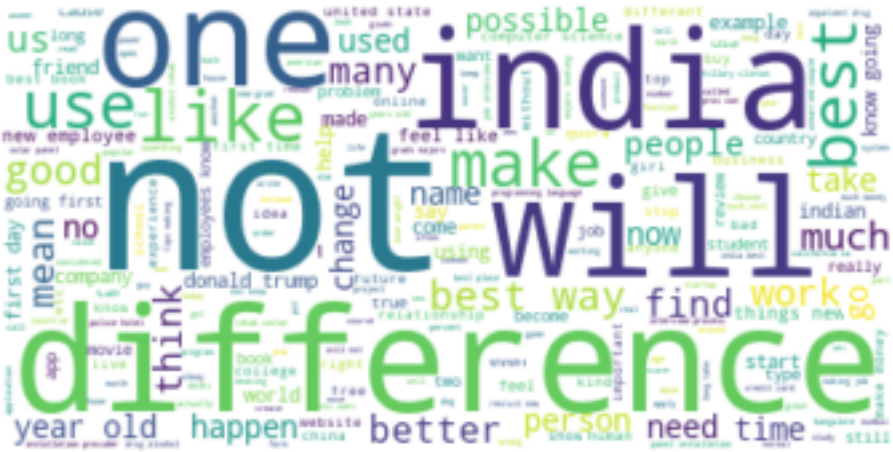

3.5.1.2 Pair plot of features ['ctc_min', 'cwc_min', 'csc_min', 'token_sort_ratio']

```
In [27]: n = df.shape[0]
         sns.pairplot(df[['ctc_min', 'cwc_min', 'csc_min', 'token_sort_ratio', 'is_duplicate']]
         plt.show()
```

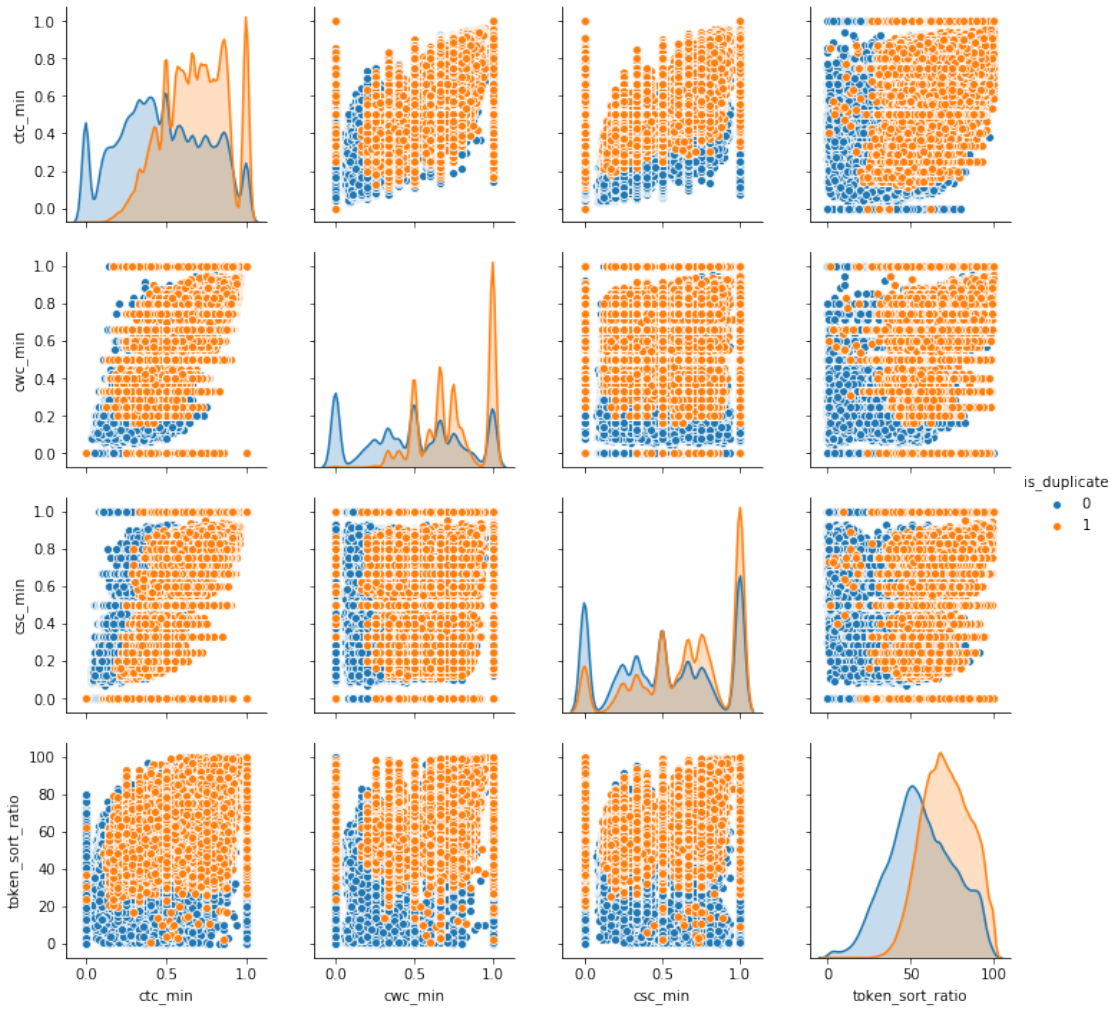

In [28]:
```
# Distribution of the token_sort_ratio
plt.figure(figsize=(10, 8))

plt.subplot(1,2,1)
sns.violinplot(x = 'is_duplicate', y = 'token_sort_ratio', data = df[0:] , )

plt.subplot(1,2,2)
sns.distplot(df[df['is_duplicate'] == 1.0]['token_sort_ratio'][0:] , label = "1", colo
sns.distplot(df[df['is_duplicate'] == 0.0]['token_sort_ratio'][0:] , label = "0" , col
plt.show()
```

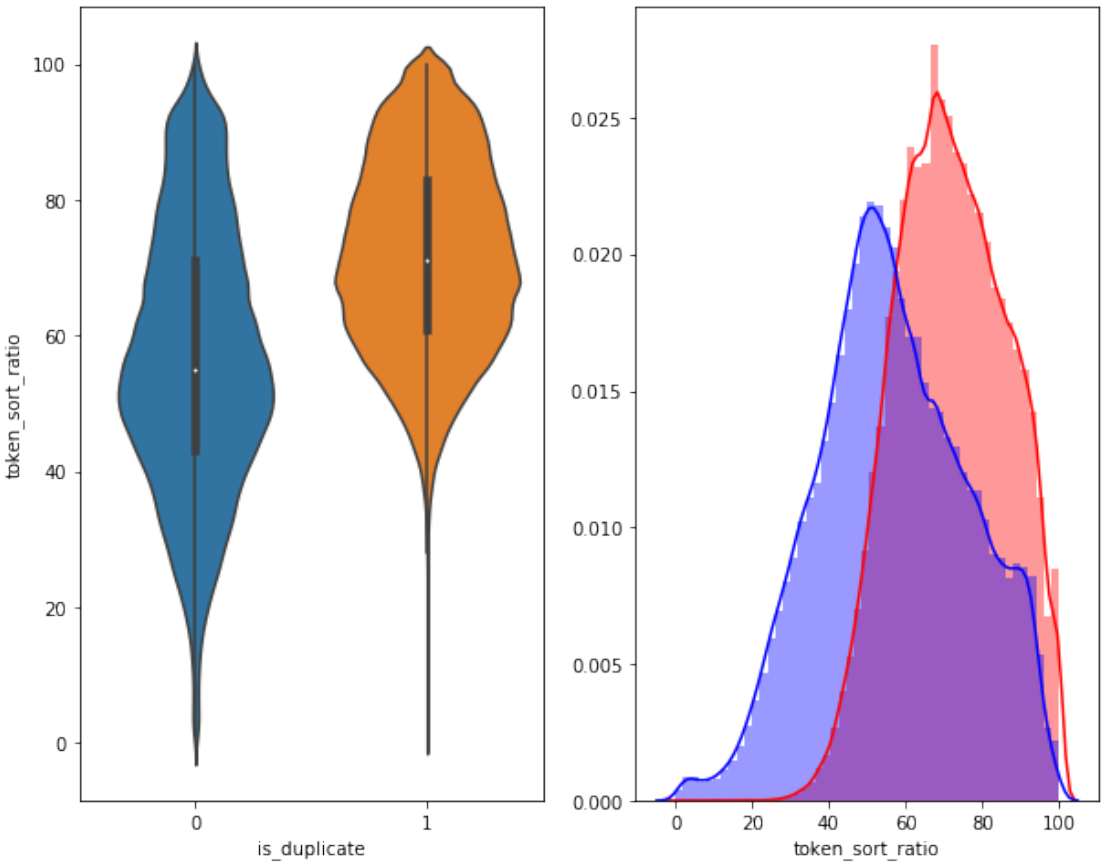


```
In [29]: plt.figure(figsize=(10, 8))

         plt.subplot(1,2,1)
         sns.violinplot(x = 'is_duplicate', y = 'fuzz_ratio', data = df[0:] , )

         plt.subplot(1,2,2)
         sns.distplot(df[df['is_duplicate'] == 1.0]['fuzz_ratio'][0:] , label = "1", color = '
         sns.distplot(df[df['is_duplicate'] == 0.0]['fuzz_ratio'][0:] , label = "0" , color =
         plt.show()
```

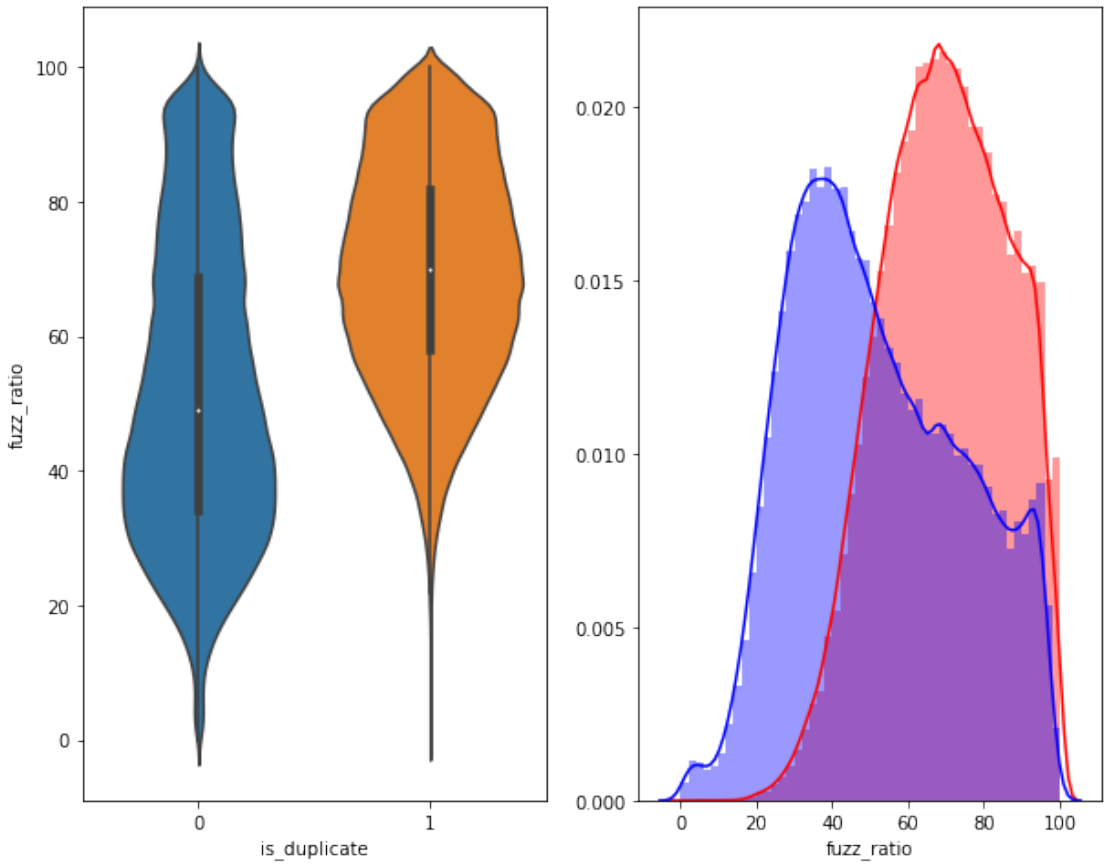

3.5.2 Visualization

```
In [30]: # Using TSNE for Dimentionality reduction for 15 Features(Generated after cleaning th
         dfp_subsampled = df[0:5000]
         X = MinMaxScaler().fit_transform(dfp_subsampled[['cwc_min', 'cwc_max', 'csc_min', 'csc
         y = dfp_subsampled['is_duplicate'].values
```

```
In [31]: tsne2d = TSNE(
             n_components=2,
             init='random', # pca
             random_state=101,
             method='barnes_hut',
             n_iter=1000,
             verbose=2,
             angle=0.5
         ).fit_transform(X)
```

```
[t-SNE] Computing 91 nearest neighbors...
[t-SNE] Indexed 5000 samples in 0.050s...
[t-SNE] Computed neighbors for 5000 samples in 1.414s...
[t-SNE] Computed conditional probabilities for sample 1000 / 5000
```

```
[t-SNE] Computed conditional probabilities for sample 2000 / 5000
[t-SNE] Computed conditional probabilities for sample 3000 / 5000
[t-SNE] Computed conditional probabilities for sample 4000 / 5000
[t-SNE] Computed conditional probabilities for sample 5000 / 5000
[t-SNE] Mean sigma: 0.130805
[t-SNE] Computed conditional probabilities in 0.544s
[t-SNE] Iteration 50: error = 81.2190323, gradient norm = 0.0482754 (50 iterations in 21.395s)
[t-SNE] Iteration 100: error = 70.6233063, gradient norm = 0.0106190 (50 iterations in 15.023s)
[t-SNE] Iteration 150: error = 68.9020691, gradient norm = 0.0049965 (50 iterations in 13.821s)
[t-SNE] Iteration 200: error = 68.1042328, gradient norm = 0.0042044 (50 iterations in 14.490s)
[t-SNE] Iteration 250: error = 67.5993195, gradient norm = 0.0047821 (50 iterations in 14.768s)
[t-SNE] KL divergence after 250 iterations with early exaggeration: 67.599319
[t-SNE] Iteration 300: error = 1.7923965, gradient norm = 0.0011835 (50 iterations in 16.494s)
[t-SNE] Iteration 350: error = 1.3952127, gradient norm = 0.0004869 (50 iterations in 16.222s)
[t-SNE] Iteration 400: error = 1.2288457, gradient norm = 0.0002799 (50 iterations in 17.308s)
[t-SNE] Iteration 450: error = 1.1391127, gradient norm = 0.0001891 (50 iterations in 17.046s)
[t-SNE] Iteration 500: error = 1.0837665, gradient norm = 0.0001422 (50 iterations in 17.321s)
[t-SNE] Iteration 550: error = 1.0480769, gradient norm = 0.0001122 (50 iterations in 15.686s)
[t-SNE] Iteration 600: error = 1.0243345, gradient norm = 0.0000994 (50 iterations in 17.401s)
[t-SNE] Iteration 650: error = 1.0078195, gradient norm = 0.0000874 (50 iterations in 17.597s)
[t-SNE] Iteration 700: error = 0.9963016, gradient norm = 0.0000808 (50 iterations in 17.444s)
[t-SNE] Iteration 750: error = 0.9881218, gradient norm = 0.0000734 (50 iterations in 16.976s)
[t-SNE] Iteration 800: error = 0.9819695, gradient norm = 0.0000697 (50 iterations in 17.552s)
[t-SNE] Iteration 850: error = 0.9767520, gradient norm = 0.0000689 (50 iterations in 17.192s)
[t-SNE] Iteration 900: error = 0.9722162, gradient norm = 0.0000615 (50 iterations in 16.717s)
[t-SNE] Iteration 950: error = 0.9683844, gradient norm = 0.0000606 (50 iterations in 16.855s)
[t-SNE] Iteration 1000: error = 0.9650941, gradient norm = 0.0000648 (50 iterations in 17.426s)
[t-SNE] Error after 1000 iterations: 0.965094
```

In [32]:
```
df = pd.DataFrame({'x':tsne2d[:,0], 'y':tsne2d[:,1] ,'label':y})

    # draw the plot in appropriate place in the grid
    sns.lmplot(data=df, x='x', y='y', hue='label', fit_reg=False, size=8,palette="Set1",ma
    plt.title("perplexity : {} and max_iter : {}".format(30, 1000))
    plt.show()
```

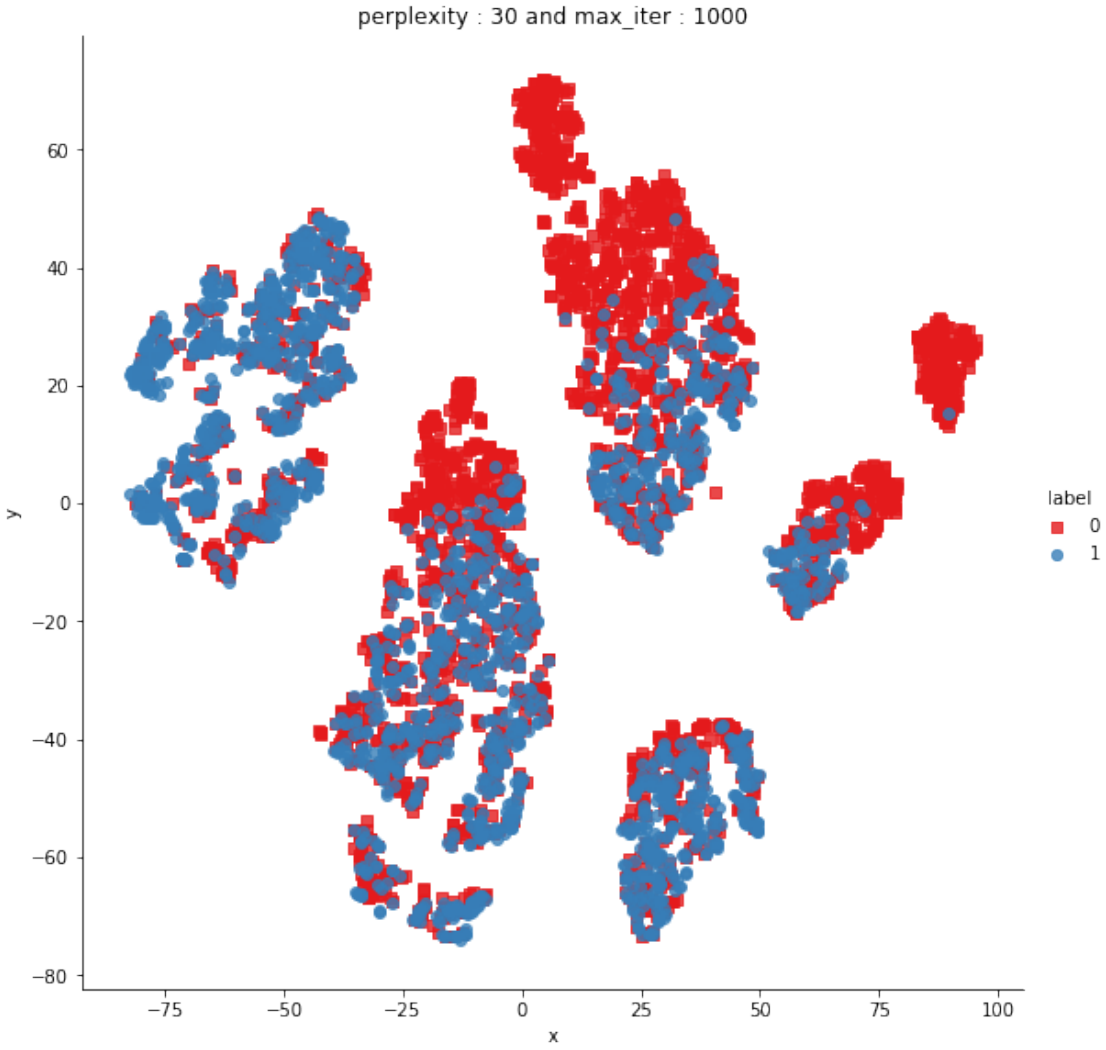


```
In [33]: tsne3d = TSNE(
             n_components=3,
             init='random', # pca
             random_state=101,
             method='barnes_hut',
             n_iter=1000,
             verbose=2,
             angle=0.5
         ).fit_transform(X)
```

```
[t-SNE] Computing 91 nearest neighbors...
[t-SNE] Indexed 5000 samples in 0.025s...
[t-SNE] Computed neighbors for 5000 samples in 1.121s...
[t-SNE] Computed conditional probabilities for sample 1000 / 5000
[t-SNE] Computed conditional probabilities for sample 2000 / 5000
```

```
[t-SNE] Computed conditional probabilities for sample 3000 / 5000
[t-SNE] Computed conditional probabilities for sample 4000 / 5000
[t-SNE] Computed conditional probabilities for sample 5000 / 5000
[t-SNE] Mean sigma: 0.130805
[t-SNE] Computed conditional probabilities in 0.539s
[t-SNE] Iteration 50: error = 80.4903412, gradient norm = 0.0297925 (50 iterations in 43.460s)
[t-SNE] Iteration 100: error = 69.3758163, gradient norm = 0.0036121 (50 iterations in 21.201s)
[t-SNE] Iteration 150: error = 67.9860153, gradient norm = 0.0016561 (50 iterations in 17.791s)
[t-SNE] Iteration 200: error = 67.4307098, gradient norm = 0.0011456 (50 iterations in 15.392s)
[t-SNE] Iteration 250: error = 67.1182632, gradient norm = 0.0008219 (50 iterations in 15.818s)
[t-SNE] KL divergence after 250 iterations with early exaggeration: 67.118263
[t-SNE] Iteration 300: error = 1.5233750, gradient norm = 0.0007065 (50 iterations in 18.442s)
[t-SNE] Iteration 350: error = 1.1845187, gradient norm = 0.0002118 (50 iterations in 24.129s)
[t-SNE] Iteration 400: error = 1.0397183, gradient norm = 0.0001055 (50 iterations in 25.252s)
[t-SNE] Iteration 450: error = 0.9664599, gradient norm = 0.0000643 (50 iterations in 26.463s)
[t-SNE] Iteration 500: error = 0.9277458, gradient norm = 0.0000571 (50 iterations in 26.338s)
[t-SNE] Iteration 550: error = 0.9077563, gradient norm = 0.0000484 (50 iterations in 22.700s)
[t-SNE] Iteration 600: error = 0.8953844, gradient norm = 0.0000408 (50 iterations in 22.267s)
[t-SNE] Iteration 650: error = 0.8859797, gradient norm = 0.0000376 (50 iterations in 25.710s)
[t-SNE] Iteration 700: error = 0.8784768, gradient norm = 0.0000307 (50 iterations in 43.087s)
[t-SNE] Iteration 750: error = 0.8713742, gradient norm = 0.0000291 (50 iterations in 38.362s)
[t-SNE] Iteration 800: error = 0.8657227, gradient norm = 0.0000288 (50 iterations in 36.017s)
[t-SNE] Iteration 850: error = 0.8613902, gradient norm = 0.0000282 (50 iterations in 37.600s)
[t-SNE] Iteration 900: error = 0.8577914, gradient norm = 0.0000244 (50 iterations in 37.989s)
[t-SNE] Iteration 950: error = 0.8543377, gradient norm = 0.0000248 (50 iterations in 25.921s)
[t-SNE] Iteration 1000: error = 0.8509857, gradient norm = 0.0000242 (50 iterations in 26.325s)
[t-SNE] Error after 1000 iterations: 0.850986
```

In [34]:
```
trace1 = go.Scatter3d(
    x=tsne3d[:,0],
    y=tsne3d[:,1],
    z=tsne3d[:,2],
    mode='markers',
    marker=dict(
        sizemode='diameter',
        color = y,
        colorscale = 'Portland',
        colorbar = dict(title = 'duplicate'),
        line=dict(color='rgb(255, 255, 255)'),
        opacity=0.75
    )
)

data=[trace1]
layout=dict(height=800, width=800, title='3d embedding with engineered features')
fig=dict(data=data, layout=layout)
py.iplot(fig, filename='3DBubble')
```