

3.Q_Mean_W2V

April 18, 2019

3.6 Featurizing text data with tfidf weighted word-vectors

```
In [42]: import pandas as pd
import matplotlib.pyplot as plt
import re
import time
import warnings
import numpy as np
from nltk.corpus import stopwords
from sklearn.preprocessing import normalize
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.feature_extraction.text import TfidfVectorizer
warnings.filterwarnings("ignore")
import sys
import os
import pandas as pd
import numpy as np
from tqdm import tqdm

# extract word2vec vectors
# https://github.com/explosion/spaCy/issues/1721
# http://landinghub.visualstudio.com/visual-cpp-build-tools
import spacy

In [43]: # avoid decoding problems
df = pd.read_csv("train.csv")

# encode questions to unicode
# https://stackoverflow.com/a/6812069
# ----- python 2 -----
# df['question1'] = df['question1'].apply(lambda x: unicode(str(x), "utf-8"))
# df['question2'] = df['question2'].apply(lambda x: unicode(str(x), "utf-8"))
# ----- python 3 -----
df['question1'] = df['question1'].apply(lambda x: str(x))
df['question2'] = df['question2'].apply(lambda x: str(x))

In [44]: df.head()
```

```

Out [44]:      id  qid1  qid2      question1 \
0    0    1    2  What is the step by step guide to invest in sh...
1    1    3    4  What is the story of Kohinoor (Koh-i-Noor) Dia...
2    2    5    6  How can I increase the speed of my internet co...
3    3    7    8  Why am I mentally very lonely? How can I solve...
4    4    9   10  Which one dissolve in water quikly sugar, salt...

      question2  is_duplicate
0  What is the step by step guide to invest in sh...      0
1  What would happen if the Indian government sto...      0
2  How can Internet speed be increased by hacking...      0
3  Find the remainder when  $23^{24}$  i...      0
4                Which fish would survive in salt water?      0

```

```

In [45]: # merge texts
questions = list(df['question1']) + list(df['question2'])

tfidf = TfidfVectorizer(lowercase=False)
tfidf.fit_transform(questions)

# dict key:word and value:tf-idf score
word2tfidf = dict(zip(tfidf.get_feature_names(), tfidf.idf_))

```

```
In [ ]: word2tfidf
```

- After we find TF-IDF scores, we convert each question to a weighted average of word2vec vectors by these scores.
- here we use a pre-trained GLOVE model which comes free with “Spacy”.
<https://spacy.io/usage/vectors-similarity>
- It is trained on Wikipedia and therefore, it is stronger in terms of word semantics.

```

In [23]: # en_vectors_web_lg, which includes over 1 million unique vectors.
nlp = spacy.load(r"C:\Users\Byron\AppData\Local\Programs\Python\PythonMaster\Lib\site-packages\en_core.

```

```

In [32]: vecs1 = []
# https://github.com/noamraph/tqdm
# tqdm is used to print the progress bar
for qu1 in tqdm(list(df['question1'])):
    doc1 = nlp(qu1)
    # 384 is the number of dimensions of vectors
    mean_vec1 = np.zeros([len(doc1), 300])
    for word1 in doc1:
        # word2vec
        vec1 = word1.vector
        # fetch df score
        try:
            idf = word2tfidf[str(word1)]
        except:
            idf = 0

```

```

        # compute final vec
        mean_vec1 += vec1 * idf
    mean_vec1 = mean_vec1.mean(axis=0)
    vecs1.append(mean_vec1)
df['q1_feats_m'] = list(vecs1)

```

100%| 404290/404290 [56:22<00:00, 119.54it/s]

```

In [33]: vecs2 = []
for qu2 in tqdm(list(df['question2'])):
    doc2 = nlp(qu2)
    mean_vec2 = np.zeros([len(doc2), 300])
    for word2 in doc2:
        # word2vec
        vec2 = word2.vector
        # fetch df score
        try:
            idf = word2tfidf[str(word2)]
        except:
            #print word
            idf = 0
        # compute final vec
        mean_vec2 += vec2 * idf
    mean_vec2 = mean_vec2.mean(axis=0)
    vecs2.append(mean_vec2)
df['q2_feats_m'] = list(vecs2)

```

100%| 404290/404290 [58:32<00:00, 115.11it/s]

```

In [34]: #prepro_features_train.csv (Simple Preprocessing Feartures)
#nlp_features_train.csv (NLP Features)
if os.path.isfile('nlp_features_train.csv'):
    dfnlp = pd.read_csv("nlp_features_train.csv",encoding='latin-1')
else:
    print("download nlp_features_train.csv from drive or run previous notebook")

if os.path.isfile('df_fe_without_preprocessing_train.csv'):
    dfppro = pd.read_csv("df_fe_without_preprocessing_train.csv",encoding='latin-1')
else:
    print("download df_fe_without_preprocessing_train.csv from drive or run previous notebook")

```

```

In [35]: df1 = dfnlp.drop(['qid1','qid2','question1','question2'],axis=1)
df2 = dfppro.drop(['qid1','qid2','question1','question2','is_duplicate'],axis=1)
df3 = df.drop(['qid1','qid2','question1','question2','is_duplicate'],axis=1)
df3_q1 = pd.DataFrame(df3.q1_feats_m.values.tolist(), index= df3.index)
df3_q2 = pd.DataFrame(df3.q2_feats_m.values.tolist(), index= df3.index)

```

```
In [36]: # dataframe of nlp features
df1.head()
```

```
Out [36]:
```

	id	is_duplicate	cwc_min	cwc_max	csc_min	csc_max	ctc_min	\
0	0	0	0.999980	0.833319	0.999983	0.999983	0.916659	
1	1	0	0.799984	0.399996	0.749981	0.599988	0.699993	
2	2	0	0.399992	0.333328	0.399992	0.249997	0.399996	
3	3	0	0.000000	0.000000	0.000000	0.000000	0.000000	
4	4	0	0.399992	0.199998	0.999950	0.666644	0.571420	

	ctc_max	last_word_eq	first_word_eq	abs_len_diff	mean_len	\
0	0.785709	0.0	1.0	2.0	13.0	
1	0.466664	0.0	1.0	5.0	12.5	
2	0.285712	0.0	1.0	4.0	12.0	
3	0.000000	0.0	0.0	2.0	12.0	
4	0.307690	0.0	1.0	6.0	10.0	

	token_set_ratio	token_sort_ratio	fuzz_ratio	fuzz_partial_ratio	\
0	100	93	93	100	
1	86	63	66	75	
2	63	63	43	47	
3	28	24	9	14	
4	67	47	35	56	

	longest_substr_ratio
0	1.000000
1	0.607843
2	0.169492
3	0.040000
4	0.179487

```
In [37]: # data before preprocessing
df2.head()
```

```
Out [37]:
```

	id	freq_qid1	freq_qid2	q1len	q2len	q1_n_words	q2_n_words	\
0	0	1	1	66	57	14	12	
1	1	4	1	51	88	8	13	
2	2	1	1	73	59	14	10	
3	3	1	1	50	65	11	9	
4	4	3	1	76	39	13	7	

	word_Common	word_Total	word_share	freq_q1+q2	freq_q1-q2
0	10.0	23.0	0.434783	2	0
1	4.0	20.0	0.200000	5	3
2	4.0	24.0	0.166667	2	0
3	0.0	19.0	0.000000	2	0
4	2.0	20.0	0.100000	4	2

```
In [38]: # Questions 1 tfidf weighted word2vec
df3_q1.head()
```

```

Out [38]:
      0      1      2      3      4      5      6  \
0 -5.856872 17.449559  4.862720  7.971019 20.345586 -5.514759 -4.077800
1  9.356103 13.098566 18.945098 -2.079594 -15.703841 -2.173409  8.969065
2  0.909520 16.050299 -8.126856 -4.848289 -2.806190  9.752280  4.349992
3 -4.950745 17.098874 -15.474965  1.044680 -2.392017 -0.051889  2.650595
4 -11.520302 19.769948 -4.510997 -6.548994 -20.835286 33.663909 -30.390504

      7      8      9      ...      290      291  \
0 -2.820742  8.029026 146.599092  ...    -17.370964  5.393082
1 -20.458267 -20.674299  13.760798  ...     25.948247  0.603713
2 -5.120332  6.785252 106.342974  ...    -20.942061  2.398984
3 -8.451192  2.584123 116.184408  ...     -2.551312 -4.971480
4  0.826553 -19.571472  84.458577  ...     -8.331733 -4.866335

      292      293      294      295      296      297      298  \
0  0.384676 -8.362788 -1.880290 -10.799672 -12.999799  3.225858  1.256145
1 -10.516349  6.040723 30.476707  3.976890 -28.254610 12.613432 -7.770673
2  8.663028 -0.654124 16.220601 -2.719094 10.485332 -1.103132 -7.290877
3 -0.478381 -1.930166  9.336016  2.574459  4.803863 -1.182989 -2.962115
4 18.828458 -40.357679 -10.336167 15.294630 -0.989347 -9.072091 -8.194567

      299
0 16.807275
1 31.456654
2 19.314250
3  3.225704
4 23.847560

[5 rows x 300 columns]

```

```

In [39]: # Questions 2 tfidf weighted word2vec
df3_q2.head()

```

```

Out [39]:
      0      1      2      3      4      5      6  \
0  0.398579 13.991607 -0.504564  9.254431 13.906436 -4.777694 -5.274421
1  4.649688  9.974928 20.330103 -0.440372 -18.128566 -1.984671  4.906458
2 -17.305105 17.355614 -9.135664 -6.038550 -1.831651  4.547895 17.935764
3  3.897911  2.545857 -2.053792  3.385450  3.424216 -2.282545 -11.763825
4 -5.391206  1.767221  1.810128 -4.097073 -3.623262  8.417368 -25.246265

      7      8      9      ...      290      291  \
0 -0.201208  4.940558 134.735950  ...    -17.810438  7.231024
1 -27.797837 -21.262646  96.965297  ...     23.015827  3.435464
2 -4.799029  3.100311  99.380095  ...    -24.310109 -1.216773
3  6.692485  5.797674  94.978085  ...     -5.435584  1.672591
4  7.473430 -2.789541  89.594627  ...    -10.407441 -8.444207

      292      293      294      295      296      297  \

```

0	1.531186	-7.528823	0.473802	-11.864658	-11.293788	1.866265
1	-5.169600	7.102491	34.516881	6.177686	-27.770856	12.926435
2	11.909693	9.591573	11.846737	1.397859	6.454157	-0.271460
3	-0.863278	-2.906553	-3.466688	-3.867892	-4.249463	-12.551012
4	-14.450059	-12.709382	-4.449050	12.563987	-11.721362	-16.459300

	298	299
0	3.616046	11.971096
1	-4.564559	33.919834
2	-12.500337	27.634567
3	4.494087	-6.223341
4	3.626297	-9.790615

[5 rows x 300 columns]

```
In [40]: print("Number of features in nlp dataframe :", df1.shape[1])
print("Number of features in preprocessed dataframe :", df2.shape[1])
print("Number of features in question1 w2v dataframe :", df3_q1.shape[1])
print("Number of features in question2 w2v dataframe :", df3_q2.shape[1])
print("Number of features in final dataframe :", df1.shape[1]+df2.shape[1]+df3_q1.sh
```

```
Number of features in nlp dataframe : 17
Number of features in preprocessed dataframe : 12
Number of features in question1 w2v dataframe : 300
Number of features in question2 w2v dataframe : 300
Number of features in final dataframe : 629
```

```
In [41]: # storing the final features to csv file
if not os.path.isfile('final_features.csv'):
    df3_q1['id']=df1['id']
    df3_q2['id']=df1['id']
    df1 = df1.merge(df2, on='id',how='left')
    df2 = df3_q1.merge(df3_q2, on='id',how='left')
    result = df1.merge(df2, on='id',how='left')
    result.to_csv('final_features.csv')
```