

## 4.ML\_models

April 18, 2019

```
In [1]: #System:
import os

#Data structures for in memory:
import csv
import pandas as pd
import numpy as np
from scipy.sparse import hstack
import math

#Database store:
from sqlalchemy import create_engine # database connection
import sqlite3

#Date and time:
import time
import datetime as dt

#Plotting:
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.manifold import TSNE

#Data transformations:
import re
from nltk.corpus import stopwords
from sklearn.feature_extraction.text import CountVectorizer, TfidfVectorizer
from sklearn.decomposition import TruncatedSVD
from sklearn.preprocessing import normalize

#Parameter tuning:
from sklearn.cross_validation import StratifiedKFold
from collections import Counter, defaultdict
from sklearn.model_selection import train_test_split, GridSearchCV, cross_val_score, RandomizedSearchCV

#Models:
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import MultinomialNB
```

```

from sklearn.naive_bayes import GaussianNB
from sklearn.linear_model import LogisticRegression, SGDClassifier
from sklearn.svm import LinearSVC
from sklearn.multiclass import OneVsRestClassifier
from sklearn.calibration import CalibratedClassifierCV
from mlxtend.classifier import StackingClassifier
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
import xgboost as xgb

#Model evaluation metrics:
from sklearn.metrics import precision_recall_curve, auc, roc_curve, confusion_matrix,

#Switch off warnings:
import warnings
warnings.filterwarnings("ignore")

from sklearn.externals import joblib
from prettytable import PrettyTable

```

```

c:\users\byron\applications\pythonmaster\lib\site-packages\sklearn\cross_validation.py:41: DeprecationWarning: This module will be removed in 0.20.", DeprecationWarning)
c:\users\byron\applications\pythonmaster\lib\site-packages\sklearn\ensemble\weight_boosting.py:271: DeprecationWarning: numpy.core.umath_tests import inner1d

```

## 4. Machine Learning Models

### 4.1 Reading data from file and storing into sql table

```

In [2]: #Creating db file from csv
if not os.path.isfile('train.db'):
    disk_engine = create_engine('sqlite:///train.db')
    start = dt.datetime.now()
    chunksize = 50000
    j = 0
    index_start = 1
    for df in pd.read_csv('final_features.csv', names=['Unnamed: 0', 'id', 'is_duplicate']):
        df.index += index_start
        j+=1
        print('{} rows'.format(j*chunksize))
        df.to_sql('data', disk_engine, if_exists='append')
        index_start = df.index[-1] + 1

In [3]: #http://www.sqlitetutorial.net/sqlite-python/create-tables/
def create_connection(db_file):
    """ create a database connection to the SQLite database
        specified by db_file
    :param db_file: database file
    :return: Connection object or None
    """

```

```

"""
try:
    conn = sqlite3.connect(db_file)
    return conn
except Error as e:
    print(e)

return None

def checkTableExists(dbcon):
    cursr = dbcon.cursor()
    str = "select name from sqlite_master where type='table'"
    table_names = cursr.execute(str)
    print("Tables in the database:")
    tables = table_names.fetchall()
    print(tables[0][0])
    return(len(tables))

```

```

In [4]: read_db = 'train.db'
        conn_r = create_connection(read_db)
        checkTableExists(conn_r)
        conn_r.close()

```

Tables in the database:  
data

```

In [5]: # try to sample data according to the computing power you have
        if os.path.isfile(read_db):
            conn_r = create_connection(read_db)
            if conn_r is not None:
                # for selecting random points
                data = pd.read_sql_query("SELECT * From data ORDER BY RANDOM() LIMIT 100001;",
                    conn_r.close()

```

```

In [6]: # remove the first row
        data.drop(data.index[0], inplace=True)
        y_true = data['is_duplicate']
        data.drop(['Unnamed: 0', 'id', 'index', 'is_duplicate'], axis=1, inplace=True)

```

```

In [7]: data.head()

```

```

Out[7]:

```

	cwc_min	cwc_max	csc_min	csc_max \
1	0.66664444518516	0.66664444518516	0.999975000624984	0.999975000624984
2	0.499987500312492	0.285710204139941	0.66664444518516	0.399992000159997
3	0.66664444518516	0.399992000159997	0.749981250468738	0.333329629670781
4	0.749981250468738	0.499991666805553	0.749981250468738	0.749981250468738
5	0.499987500312492	0.19999800002	0.999980000399992	0.454541322351615

	ctc_min	ctc_max	last_word_eq	first_word_eq	\
1	0.857130612419823	0.857130612419823	0.0	1.0	
2	0.571420408279882	0.307689940846609	0.0	1.0	
3	0.714275510349852	0.312498046887207	1.0	0.0	
4	0.749990625117186	0.599994000059999	1.0	1.0	
5	0.77776913589849	0.259258299043337	1.0	0.0	

	abs_len_diff	mean_len	...	290_y	\
1	0.0	7.0	...	-13.684094414115	
2	6.0	10.0	...	-2.89921551942825	
3	9.0	11.5	...	-32.1513776183128	
4	2.0	9.0	...	1.11949726939201	
5	18.0	18.0	...	-33.0604563355446	

	291_y	292_y	293_y	294_y	\
1	5.18617536127567	3.69274061173201	-1.06995718181133	-3.38792563974857	
2	2.25030846148729	-4.55699910968542	3.22107343003154	4.69975774548948	
3	-6.48564624227583	-10.6156985536218	-8.61111462116241	-2.86723747849464	
4	-5.4107170291245	-3.75332200527191	-4.40629441710189	0.106489285826683	
5	30.6366586647928	10.6500630229712	-10.6027148663998	8.85900411009789	

	295_y	296_y	297_y	\
1	1.86694558337331	-1.14174094796181	-3.97690352797508	
2	1.60277144983411	-4.12365251034498	-7.01728013157845	
3	15.2251101061702	-9.93901033699512	-3.61792010068893	
4	3.812221378088	-5.15402545034885	-0.789197444915772	
5	-24.4780361577868	-4.69883567839861	-13.0958931222558	

	298_y	299_y
1	-2.48735983669758	4.12184119224548
2	0.270279049873352	8.08513672836125
3	-3.56169393658638	-3.51276577170938
4	1.25951708108187	3.11145649943501
5	26.8136106580496	-25.8301425874233

[5 rows x 626 columns]

In [8]: data.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 100000 entries, 1 to 100000
Columns: 626 entries, cwc_min to 299_y
dtypes: object(626)
memory usage: 478.4+ MB
```

In [9]: col = data.columns.values

## 4.2 Converting strings to numerics

```
In [10]: data = pd.DataFrame(np.array(data).astype(float), columns = col)

In [11]: # https://stackoverflow.com/questions/7368789/convert-all-strings-in-a-list-to-int
y_true = list(map(int, y_true.values))
```

## 4.3 Random train test split( 70:30)

```
In [12]: X_train, X_test, y_train, y_test = train_test_split(data, y_true, stratify=y_true, test_size=0.3)

In [13]: print("Number of data points in train data :", X_train.shape)
          print("Number of data points in test data :", X_test.shape)
```

Number of data points in train data : (70000, 626)

Number of data points in test data : (30000, 626)

```
In [14]: print("-"*10, "Distribution of output variable in train data", "-"*10)
          train_distr = Counter(y_train)
          train_len = len(y_train)
          print("Class 0: ", int(train_distr[0])/train_len, "Class 1: ", int(train_distr[1])/train_len)
          print("-"*10, "Distribution of output variable in test data", "-"*10)
          test_distr = Counter(y_test)
          test_len = len(y_test)
          print("Class 0: ", int(test_distr[0])/test_len, "Class 1: ", int(test_distr[1])/test_len)
```

----- Distribution of output variable in train data -----

Class 0: 0.6283285714285715 Class 1: 0.3716714285714286

----- Distribution of output variable in test data -----

Class 0: 0.37166666666666665 Class 1: 0.37166666666666665

```
In [15]: # This function plots the confusion matrices given y_i, y_i_hat.
def plot_confusion_matrix(test_y, predict_y):
    C = confusion_matrix(test_y, predict_y)
    # C = 9,9 matrix, each cell (i,j) represents number of points of class i are predicted as class j

    A = (((C.T)/(C.sum(axis=1))).T)
    #divid each element of the confusion matrix with the sum of elements in that column

    # C = [[1, 2],
    #      [3, 4]]
    # C.T = [[1, 3],
    #        [2, 4]]
    # C.sum(axis = 1) axis=0 corresponds to columns and axis=1 corresponds to rows in C
    # C.sum(axis = 1) = [[3, 7]]
    # ((C.T)/(C.sum(axis=1))) = [[1/3, 3/7],
    #                             [2/3, 4/7]]
```

```

# ((C.T)/(C.sum(axis=1))).T = [[1/3, 2/3]
#                               [3/7, 4/7]]
# sum of row elements = 1

B = (C/C.sum(axis=0))
#divid each element of the confusion matrix with the sum of elements in that row
# C = [[1, 2],
#       [3, 4]]
# C.sum(axis = 0) axis=0 corresponds to columns and axis=1 corresponds to rows in
# C.sum(axis=0) = [[4, 6]]
# (C/C.sum(axis=0)) = [[1/4, 2/6],
#                       [3/4, 4/6]]
plt.figure(figsize=(20,4))

labels = [1,2]
# representing A in heatmap format
cmap=sns.light_palette("blue")
plt.subplot(1, 3, 1)
sns.heatmap(C, annot=True, cmap=cmap, fmt=".3f", xticklabels=labels, yticklabels=labels)
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')
plt.title("Confusion matrix")

plt.subplot(1, 3, 2)
sns.heatmap(B, annot=True, cmap=cmap, fmt=".3f", xticklabels=labels, yticklabels=labels)
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')
plt.title("Precision matrix")

plt.subplot(1, 3, 3)
# representing B in heatmap format
sns.heatmap(A, annot=True, cmap=cmap, fmt=".3f", xticklabels=labels, yticklabels=labels)
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')
plt.title("Recall matrix")

plt.show()

```

#### 4.4 Building a random model (Finding worst-case log-loss)

```

In [16]: # we need to generate 9 numbers and the sum of numbers should be 1
# one solution is to generate 9 numbers and divide each of the numbers by their sum
# ref: https://stackoverflow.com/a/18662466/4084039
# we create a output array that has exactly same size as the CV data
predicted_y = np.zeros((test_len,2))
for i in range(test_len):
    rand_probs = np.random.rand(1,2)

```


```

predicted_y[i] = ((rand_probs/sum(sum(rand_probs))))[0])
print("Log loss on Test Data using Random Model",log_loss(y_test, predicted_y, eps=1e-15))

predicted_y =np.argmax(predicted_y, axis=1)
plot_confusion_matrix(y_test, predicted_y)

```

Log loss on Test Data using Random Model 0.885340305050038



#### 4.4 Logistic Regression with hyperparameter tuning

```

In [17]: alpha = [10 ** x for x in range(-5, 2)] # hyperparam for SGD classifier.

# read more about SGDClassifier() at http://scikit-learn.org/stable/modules/generated/
# -----
# default parameters
# SGDClassifier(loss=hinge, penalty=l2, alpha=0.0001, l1_ratio=0.15, fit_intercept=True,
# shuffle=True, verbose=0, epsilon=0.1, n_jobs=1, random_state=None, learning_rate=optimal,
# class_weight=None, warm_start=False, average=False, n_iter=None)

# some of methods
# fit(X, y[, coef_init, intercept_init, ])          Fit linear model with Stochastic Gradient Descent
# predict(X)          Predict class labels for samples in X.

#-----
# video link:
#-----

log_error_array=[]
for i in alpha:
    clf = SGDClassifier(alpha=i, penalty='l2', loss='log', random_state=42)
    clf.fit(X_train, y_train)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(X_train, y_train)
    predict_y = sig_clf.predict_proba(X_test)
    log_error_array.append(log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15))
    print('For values of alpha = ', i, "The log loss is:",log_loss(y_test, predict_y,

```

```

fig, ax = plt.subplots()
ax.plot(alpha, log_error_array,c='g')
for i, txt in enumerate(np.round(log_error_array,3)):
    ax.annotate((alpha[i],np.round(txt,3)), (alpha[i],log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()

best_alpha = np.argmin(log_error_array)
clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='log', random_state=42)
clf.fit(X_train, y_train)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(X_train, y_train)

predict_y = sig_clf.predict_proba(X_train)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:",log_loss(y_train,predict_y))
predict_y = sig_clf.predict_proba(X_test)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:",log_loss(y_test,predict_y))
predicted_y =np.argmax(predict_y,axis=1)
print("Total number of data points :", len(predicted_y))
plot_confusion_matrix(y_test, predicted_y)


```

```


For values of alpha = 1e-05 The log loss is: 0.48771178018319705
For values of alpha = 0.0001 The log loss is: 0.4834804577479922
For values of alpha = 0.001 The log loss is: 0.4860223659213464
For values of alpha = 0.01 The log loss is: 0.48074674780307397
For values of alpha = 0.1 The log loss is: 0.4781428484211463
For values of alpha = 1 The log loss is: 0.472357294213177
For values of alpha = 10 The log loss is: 0.5182490126857469

```





For values of best alpha = 1 The train log loss is: 0.46529034161165644  
 For values of best alpha = 1 The test log loss is: 0.472357294213177  
 Total number of data points : 30000



#### 4.5 Linear SVM with hyperparameter tuning

In [18]: `alpha = [10 ** x for x in range(-5, 2)] # hyperparam for SGD classifier.`

```
# read more about SGDClassifier() at http://scikit-learn.org/stable/modules/generated/
# -----
# default parameters
# SGDClassifier(loss=hinge, penalty=l2, alpha=0.0001, l1_ratio=0.15, fit_intercept=True,
# shuffle=True, verbose=0, epsilon=0.1, n_jobs=1, random_state=None, learning_rate=op
```

```

# class_weight=None, warm_start=False, average=False, n_iter=None)

# some of methods
# fit(X, y[, coef_init, intercept_init, ])          Fit linear model with Stochastic Gradient Descent
# predict(X)          Predict class labels for samples in X.

#-----
# video link:
#-----

log_error_array=[]
for i in alpha:
    clf = SGDClassifier(alpha=i, penalty='l1', loss='hinge', random_state=42)
    clf.fit(X_train, y_train)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(X_train, y_train)
    predict_y = sig_clf.predict_proba(X_test)
    log_error_array.append(log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15))
    print('For values of alpha = ', i, "The log loss is:", log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15))

fig, ax = plt.subplots()
ax.plot(alpha, log_error_array, c='g')
for i, txt in enumerate(np.round(log_error_array, 3)):
    ax.annotate((alpha[i], np.round(txt, 3)), (alpha[i], log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()

best_alpha = np.argmin(log_error_array)
clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l1', loss='hinge', random_state=42)
clf.fit(X_train, y_train)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(X_train, y_train)

predict_y = sig_clf.predict_proba(X_train)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:", log_loss(y_train, predict_y, labels=clf.classes_, eps=1e-15))
predict_y = sig_clf.predict_proba(X_test)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:", log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15))
predicted_y = np.argmax(predict_y, axis=1)
print("Total number of data points :", len(predicted_y))
plot_confusion_matrix(y_test, predicted_y)


```

```


For values of alpha = 1e-05 The log loss is: 0.4539013938818281
For values of alpha = 0.0001 The log loss is: 0.4582849827404135

```

For values of alpha = 0.001 The log loss is: 0.47961404331751945  
 For values of alpha = 0.01 The log loss is: 0.5372885487844338  
 For values of alpha = 0.1 The log loss is: 0.5682885200407735  
 For values of alpha = 1 The log loss is: 0.6289676976712478  
 For values of alpha = 10 The log loss is: 0.6448495605353364



For values of best alpha = 1e-05 The train log loss is: 0.4485062220124629  
 For values of best alpha = 1e-05 The test log loss is: 0.4539013938818281  
 Total number of data points : 30000



## 4.6 XGBoost

```

In [19]: params = {}
         params['objective'] = 'binary:logistic'
         params['eval_metric'] = 'logloss'
         params['eta'] = 0.02
         params['max_depth'] = 4

         d_train = xgb.DMatrix(X_train, label=y_train)
         d_test = xgb.DMatrix(X_test, label=y_test)

         watchlist = [(d_train, 'train'), (d_test, 'valid')]

         bst = xgb.train(params, d_train, 400, watchlist, early_stopping_rounds=20, verbose_eval=10)

         xgdmatrix = xgb.DMatrix(X_train, y_train)
         predict_y = bst.predict(d_test)
         print("The test log loss is:", log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-08))

[10:31:16] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error:
[0]      train-logloss:0.684597      valid-logloss:0.684678
Multiple eval metrics have been passed: 'valid-logloss' will be used for early stopping.

Will train until valid-logloss hasn't improved in 20 rounds.
[10:31:18] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error:
[10:31:19] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error:
[10:31:20] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error:
[10:31:21] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error:
[10:31:22] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error:
[10:31:23] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error:
[10:31:24] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error:
[10:31:25] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error:
[10:31:26] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error:
[10:31:27] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error:
[10]      train-logloss:0.614357      valid-logloss:0.615202
[10:31:28] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error:
[10:31:29] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error:
[10:31:30] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error:
[10:31:31] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error:
[10:31:32] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error:
[10:31:33] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error:
[10:31:34] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error:
[10:31:35] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error:
[10:31:36] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error:
[10:31:37] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error:
[20]      train-logloss:0.562945      valid-logloss:0.564105
[10:31:38] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error:
[10:31:40] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error:
[10:31:41] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error:
[10:31:42] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error:

```



[illegible]









[illegible]





```

[10:38:41] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[10:38:42] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[10:38:43] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[10:38:44] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[10:38:46] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[10:38:47] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[10:38:48] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[380]      train-logloss:0.344861      valid-logloss:0.354599
[10:38:49] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[10:38:51] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[10:38:52] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[10:38:53] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[10:38:54] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[10:38:56] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[10:38:57] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[10:38:58] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[10:38:59] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[10:39:01] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[390]      train-logloss:0.343845      valid-logloss:0.353868
[10:39:02] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[10:39:03] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[10:39:04] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[10:39:06] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[10:39:07] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[10:39:08] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[10:39:09] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[10:39:10] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[10:39:12] C:\Users\Administrator\Desktop\xgboost\src\tree\updater_prune.cc:74: tree pruning error
[399]      train-logloss:0.342943      valid-logloss:0.353267
The test log loss is: 0.35326659471739547


```

```

In [20]: predicted_y = np.array(predict_y>0.5,dtype=int)
        print("Total number of data points :", len(predicted_y))
        plot_confusion_matrix(y_test, predicted_y)

```

Total number of data points : 30000



## 5. Assignments

1. Try out models (Logistic regression, Linear-SVM) with simple TF-IDF vectors instead of TF-IDF weighted word2Vec.
2. Hyperparameter tune XgBoost using RandomSearch to reduce the log-loss.

```
In [21]: del X_train, X_test, y_train, y_test, data, y_true
```

```
In [22]: #Read in the training data
df = pd.read_csv("train.csv")
df['question1'] = df['question1'].apply(lambda x: str(x))
df['question2'] = df['question2'].apply(lambda x: str(x))
```

```
In [23]: df = df.sample(n=100000)
```

```
In [24]: df.head(5)
```

```
Out[24]:
```

	id	qid1	qid2	\	question1	\	question2	is_duplicate
52066	52066	92247	92248		Why allow refugees in Europe when most of th...		Why should we accept refugees in Europe?	0
348520	348520	85073	477085		How is CAT percentile calculated?		How would one explain the percentile system in...	0
94223	94223	157392	157393		What does thanoo, vanno and kooi means in mala...		What does the Malayalam word AYYO mean?	0
235267	235267	345946	345947		Which planet in our solar system is the most h...		To which planet in our Solar System would you ...	0
278508	278508	397794	397795		What is a good free C compiler for Windows 7?		Where can I download a free Turbo C compiler?	0

```
In [25]: data = {'id':df['id'], 'text':df['question1'] + ' ' + df['question2'], 'is_duplicate'
df = pd.DataFrame(data=data,index=data['id'])
df.drop(labels=['id'],axis=1,inplace=True)
df.head()
```

```
Out[25]:
```

	id	text	is_duplicate
52066	52066	Why allow refugees in Europe when most of th...	0
348520	348520	How is CAT percentile calculated? How would on...	0
94223	94223	What does thanoo, vanno and kooi means in mala...	0
235267	235267	Which planet in our solar system is the most h...	0
278508	278508	What is a good free C compiler for Windows 7? ...	0

```

In [26]: del data

In [27]: X_train, X_test = train_test_split(df, stratify=df['is_duplicate'], test_size=0.3)

In [28]: X_train.shape

Out[28]: (70000, 2)

In [29]: Y_train = X_train['is_duplicate']
          X_train.drop(labels=['is_duplicate'], axis=1, inplace=True)

In [30]: X_train.shape

Out[30]: (70000, 1)

In [31]: X_test.shape

Out[31]: (30000, 2)

In [32]: Y_test = X_test['is_duplicate']
          X_test.drop(labels=['is_duplicate'], axis=1, inplace=True)

In [33]: X_test.shape

Out[33]: (30000, 1)

In [34]: tfidf = TfidfVectorizer(lowercase=True, stop_words='english', ngram_range=(1,3), use_idf=

In [35]: #TRAIN
          X_tfidf_train = tfidf.fit_transform(X_train['text'])
          X_tfidf_train = pd.DataFrame(data=X_tfidf_train.toarray(), index=X_train.index.values,
          X_tfidf_train.head()

Out[35]:
          000          10  100  1000  1000 notes  1000 rupee  \
264395  0.000000  0.000000  0.0   0.0         0.0         0.0
247334  0.000000  0.000000  0.0   0.0         0.0         0.0
162234  0.655974  0.277419  0.0   0.0         0.0         0.0
265526  0.000000  0.000000  0.0   0.0         0.0         0.0
200974  0.000000  0.000000  0.0   0.0         0.0         0.0

          1000 rupee notes  1000 rupees  1000 rupees notes  11  ...  year old  \
264395          0.0         0.0         0.0  0.0  ...         0.0
247334          0.0         0.0         0.0  0.0  ...         0.0
162234          0.0         0.0         0.0  0.0  ...         0.0
265526          0.0         0.0         0.0  0.0  ...         0.0
200974          0.0         0.0         0.0  0.0  ...         0.0

          year resolution  year resolutions  years  years old  yes  york  young  \
264395          0.0         0.0         0.0         0.0  0.0  0.0  0.0
247334          0.0         0.0         0.0         0.0  0.0  0.0  0.0

```

162234	0.0	0.0	0.0	0.0	0.0	0.0	0.0
265526	0.0	0.0	0.0	0.0	0.0	0.0	0.0
200974	0.0	0.0	0.0	0.0	0.0	0.0	0.0

	youtube	zero
264395	0.0	0.0
247334	0.0	0.0
162234	0.0	0.0
265526	0.0	0.0
200974	0.0	0.0

[5 rows x 2000 columns]

In [36]: X\_tfidf\_train.shape

Out[36]: (70000, 2000)

In [37]: #TEST

```
X_tfidf_test = tfidf.transform(X_test['text'])
X_tfidf_test = pd.DataFrame(data=X_tfidf_test.toarray(), index=X_test.index.values, columns=X_test.columns)
X_tfidf_test.head()
```

Out[37]:

	000	10	100	1000	1000	notes	1000	rupee	1000	rupee	notes	\
322177	0.0	0.0	0.0	0.0		0.0		0.0			0.0	
147359	0.0	0.0	0.0	0.0		0.0		0.0			0.0	
121243	0.0	0.0	0.0	0.0		0.0		0.0			0.0	
360595	0.0	0.0	0.0	0.0		0.0		0.0			0.0	
261909	0.0	0.0	0.0	0.0		0.0		0.0			0.0	

	1000	rupees	1000	rupees	notes	11	...	year	old	year	resolution	\
322177		0.0			0.0	0.0	...		0.0		0.0	
147359		0.0			0.0	0.0	...		0.0		0.0	
121243		0.0			0.0	0.0	...		0.0		0.0	
360595		0.0			0.0	0.0	...		0.0		0.0	
261909		0.0			0.0	0.0	...		0.0		0.0	

	year	resolutions	years	years	old	yes	york	young	youtube	zero
322177		0.0	0.0		0.0	0.0	0.0	0.0	0.0	0.0
147359		0.0	0.0		0.0	0.0	0.0	0.0	0.0	0.0
121243		0.0	0.0		0.0	0.0	0.0	0.0	0.0	0.0
360595		0.0	0.0		0.0	0.0	0.0	0.0	0.0	0.0
261909		0.0	0.0		0.0	0.0	0.0	0.0	0.0	0.0

[5 rows x 2000 columns]

In [38]: X\_tfidf\_test.shape

Out[38]: (30000, 2000)



```
In [39]: if os.path.isfile('nlp_features_train.csv'):
        dfnlp = pd.read_csv("nlp_features_train.csv",encoding='latin-1')
    else:
        print("download nlp_features_train.csv from drive or run previous notebook")

    if os.path.isfile('df_fe_without_preprocessing_train.csv'):
        dfppro = pd.read_csv("df_fe_without_preprocessing_train.csv",encoding='latin-1')
    else:
        print("download df_fe_without_preprocessing_train.csv from drive or run previous notebook")
```

```
In [40]: dfnlp.head(n=2)
```

```
Out[40]:
```

	id	qid1	qid2	question1	question2	is_duplicate	cwc_min	cwc_max	csc_min	csc_max	ctc_max	last_word_eq	abs_len_diff	mean_len	token_set_ratio	token_sort_ratio	fuzz_ratio	fuzz_partial_ratio	longest_substr_ratio
0	0	1	2	what is the step by step guide to invest in sh...	what is the step by step guide to invest in sh...	0	0.999980	0.833319	0.999983	0.999983	0.785709	0.0	1.0	2.0	13.0	100	93	100	1.000000
1	1	3	4	what is the story of kohinoor koh i noor dia...	what would happen if the indian government sto...	0	0.799984	0.399996	0.749981	0.599988	0.466664	0.0	1.0	5.0	12.5	86	63	75	0.607843

[2 rows x 21 columns]

```
In [41]: dfppro.head(n=2)
```

```
Out[41]:
```

	id	qid1	qid2	question1	question2	is_duplicate	freq_qid1	freq_qid2	q1len	q2len	q1_n_words	q2_n_words	word_Common	word_Total	word_share	freq_q1+q2	freq_q1-q2
0	0	1	2	What is the step by step guide to invest in sh...	What is the step by step guide to invest in sh...	0	1	1	66	57	14	12	10.0	23.0			
1	1	3	4	What is the story of Kohinoor (Koh-i-Noor) Dia...	What would happen if the Indian government sto...	0	4	1	51	88	8	13	4.0	20.0			

0	0.434783	2	0
1	0.200000	5	3

```
In [42]: dfnlp.drop(labels=['id','qid1','qid2','question1','question2','is_duplicate'],axis=1,
dfppro.drop(labels=['id','qid1','qid2','question1','question2','is_duplicate'],axis=1)

In [43]: X_train = np.array(X_tfidf_train.merge(dfnlp, how='inner', left_index=True, right_index=True))
Y_train = np.array(Y_train)

In [44]: X_test = np.array(X_tfidf_test.merge(dfnlp, how='inner', left_index=True, right_index=True))
Y_test = np.array(Y_test)

In [45]: del X_tfidf_train, X_tfidf_test, dfnlp, dfppro


In [64]: C = [10 ** x for x in range(-5, 2)]
log_error_array=list()
for i in C:
    clf = LogisticRegression(penalty='l2',C=i,random_state=42,n_jobs=-1)
    clf.fit(X_train, Y_train)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(X_train, Y_train)
    predict_y = sig_clf.predict_proba(X_test)
    log_error_array.append(log_loss(Y_test, predict_y, labels=clf.classes_, eps=1e-15))
    print('For values of C = ', i, "The log loss is:",log_loss(Y_test, predict_y, labels=clf.classes_, eps=1e-15))

fig, ax = plt.subplots()
ax.plot(C, log_error_array,c='g')
for i, txt in enumerate(np.round(log_error_array,3)):
    ax.annotate((C[i],np.round(txt,3)), (C[i],log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each C")
plt.xlabel("C i's")
plt.ylabel("Error measure")
plt.show()


best_C = np.argmin(log_error_array)
clf = LogisticRegression(penalty='l2',C=C[best_C],random_state=42,n_jobs=-1)
clf.fit(X_train, Y_train)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(X_train, Y_train)

predict_prob_train = sig_clf.predict_proba(X_train)
print('For values of best C = ', C[best_C], "The train log loss is:",log_loss(Y_train, predict_prob_train, labels=clf.classes_, eps=1e-15))
predict_prob_test = sig_clf.predict_proba(X_test)
print('For values of best C = ', C[best_C], "The test log loss is:",log_loss(Y_test, predict_prob_test, labels=clf.classes_, eps=1e-15))
predicted_class_test = sig_clf.predict(X_test)
print("Total number of data points :", len(predicted_class_test))
plot_confusion_matrix(Y_test, predicted_class_test)
```

For values of  $C = 1e-05$  The log loss is: 0.5082266785082132  
 For values of  $C = 0.0001$  The log loss is: 0.4670614577299387  
 For values of  $C = 0.001$  The log loss is: 0.44647625434776567  
 For values of  $C = 0.01$  The log loss is: 0.4201115739826379  
 For values of  $C = 0.1$  The log loss is: 0.3904853215500971  
 For values of  $C = 1$  The log loss is: 0.37886642262304765  
 For values of  $C = 10$  The log loss is: 0.38025203183389183



For values of best  $C = 1$  The train log loss is: 0.3634400307826507  
 For values of best  $C = 1$  The test log loss is: 0.37886642262304765  
 Total number of data points : 30000



```
In [65]: miss_class = 1 - accuracy_score(Y_test,predicted_class_test)
        print("Number of missclassified points :",round(miss_class*100,2),'%')
```

Number of missclassified points : 18.65 %

```
In [66]: auc_score = roc_auc_score(Y_test,predict_prob_test[:,1])
        print("AUC score :",round(auc_score*100,2))
```

AUC score : 90.06

```
In [68]: fpr_LR,tpr_LR,thresholds_LR = roc_curve(Y_test, predict_prob_test[:,1])
```


```
In [69]: C = [10 ** x for x in range(-5, 2)]
        log_error_array=list()
        for i in C:
            clf = LinearSVC(penalty='l2', loss='hinge', C=i, random_state=42)
            clf.fit(X_train, Y_train)
            sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
            sig_clf.fit(X_train, Y_train)
            predict_y = sig_clf.predict_proba(X_test)
            log_error_array.append(log_loss(Y_test, predict_y, labels=clf.classes_, eps=1e-15))
        print('For values of C = ', i, "The log loss is:",log_loss(Y_test, predict_y, labels=clf.classes_, eps=1e-15))
```

```
fig, ax = plt.subplots()
ax.plot(C, log_error_array,c='g')
for i, txt in enumerate(np.round(log_error_array,3)):
    ax.annotate((C[i],np.round(txt,3)), (C[i],log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each C")
plt.xlabel("C i's")
plt.ylabel("Error measure")
plt.show()
```


```
best_C = np.argmin(log_error_array)
clf = LinearSVC(penalty='l2', loss='hinge', C= C[best_C],random_state=42)
clf.fit(X_train, Y_train)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(X_train, Y_train)
```

```
predict_prob_train = sig_clf.predict_proba(X_train)
print('For values of best C = ', C[best_C], "The train log loss is:",log_loss(Y_train, predict_prob_train))
predict_prob_test = sig_clf.predict_proba(X_test)
print('For values of best C = ', C[best_C], "The test log loss is:",log_loss(Y_test, predict_prob_test))
predicted_class_test = sig_clf.predict(X_test)
print("Total number of data points :", len(predicted_class_test))
plot_confusion_matrix(Y_test, predicted_class_test)
```

For values of  $C = 1e-05$  The log loss is: 0.5012534911437849  
 For values of  $C = 0.0001$  The log loss is: 0.4692509702920933  
 For values of  $C = 0.001$  The log loss is: 0.4481429451633474  
 For values of  $C = 0.01$  The log loss is: 0.4145877516468119  
 For values of  $C = 0.1$  The log loss is: 0.4019437533693079  
 For values of  $C = 1$  The log loss is: 0.3952581358024728  
 For values of  $C = 10$  The log loss is: 0.3952581358024728



For values of best  $C = 1$  The train log loss is: 0.38663769406321974  
 For values of best  $C = 1$  The test log loss is: 0.3952581358024728  
 Total number of data points : 30000



```
In [70]: miss_class = 1 - accuracy_score(Y_test,predicted_class_test)
         print("Number of missclassified points :",round(miss_class*100,2),'%')
```

Number of missclassified points : 18.89 %

```
In [71]: auc_score = roc_auc_score(Y_test,predict_prob_test[:,1])
         print("AUC score :",round(auc_score*100,2))
```

AUC score : 89.41

```
In [72]: fpr_SVM, tpr_SVM, thresholds_SVM = roc_curve(Y_test, predict_prob_test[:,1])
```

```
In [73]: # Sklearn version:
         # n_estimators = [3,5,20,50,80,100]
         # max_depth = [3,5,7]
         # log_error_array=list()
         # for estimator in n_estimators:
         #     for depth in max_depth:
         #         clf = GradientBoostingClassifier(n_estimators=estimator, max_depth=depth, r
         #         clf.fit(X_train, Y_train)
         #         sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
         #         sig_clf.fit(X_train, Y_train)
         #         predict_y = sig_clf.predict_proba(X_test)
         #         log_error_array.append(log_loss(Y_test, predict_y, labels=clf.classes_, eps
         #         print('For values of n_estimators = ', estimator, ' and max_depth = ', dept
```

```
In [74]: # Xgboost - sklearn wrapper version:
```

```
         # Create the parameter grid: param_grid
         # model.get_params()
         param_grid = {
             'base_estimator__learning_rate': np.arange(0.05,1.05,0.05),
             'base_estimator__n_estimators': [5,10,20,40,60,80,100,120],
             'base_estimator__subsample': np.arange(0.05,1.05,0.05),
             'base_estimator__max_depth': [3,5,7,9]
         }

         if os.path.isfile('xgbClassifier.pkl') == False:
             # Model:
             base_model = xgb.XGBClassifier()
             model = CalibratedClassifierCV(base_estimator = base_model)

             # Perform random search:
             grid_search = RandomizedSearchCV(estimator=model, param_distributions=param_grid,
             grid_result = grid_search.fit(X_train,Y_train)

             # summarize results
```

```

print("Best: %f using %s" % (grid_result.best_score_, grid_result.best_params_))
means = grid_result.cv_results_['mean_test_score']
stds = grid_result.cv_results_['std_test_score']
params = grid_result.cv_results_['params']
for mean, stdev, param in zip(means, stds, params):
    print("%f (%f) with: %r" % (mean, stdev, param))

joblib.dump(grid_result.best_estimator_, 'xgbClassifier.pkl')
model = grid_result.best_estimator_
else:
    model = joblib.load('xgbClassifier.pkl')

```

```

In [75]: predict_prob_train = model.predict_proba(X_train)
print("The train log loss is:", log_loss(Y_train, predict_prob_train))
predict_prob_test = model.predict_proba(X_test)
print("The test log loss is:", log_loss(Y_test, predict_prob_test))
predicted_class_test = model.predict(X_test)

```

The train log loss is: 0.3549992799848697


The test log loss is: 0.35377406307865034

The xgbclassifier model did really well on both train and test data which implies that model is stable and does not over or under fit.

```

In [76]: plot_confusion_matrix(Y_test, predicted_class_test)

```



```

In [77]: miss_class = 1 - accuracy_score(Y_test, predicted_class_test)
print("Number of missclassified points :", round(miss_class*100,2), '%')

```

Number of missclassified points : 16.96 %

```

In [78]: auc_score = roc_auc_score(Y_test, predict_prob_test[:,1])
print("AUC score :", round(auc_score*100,2))

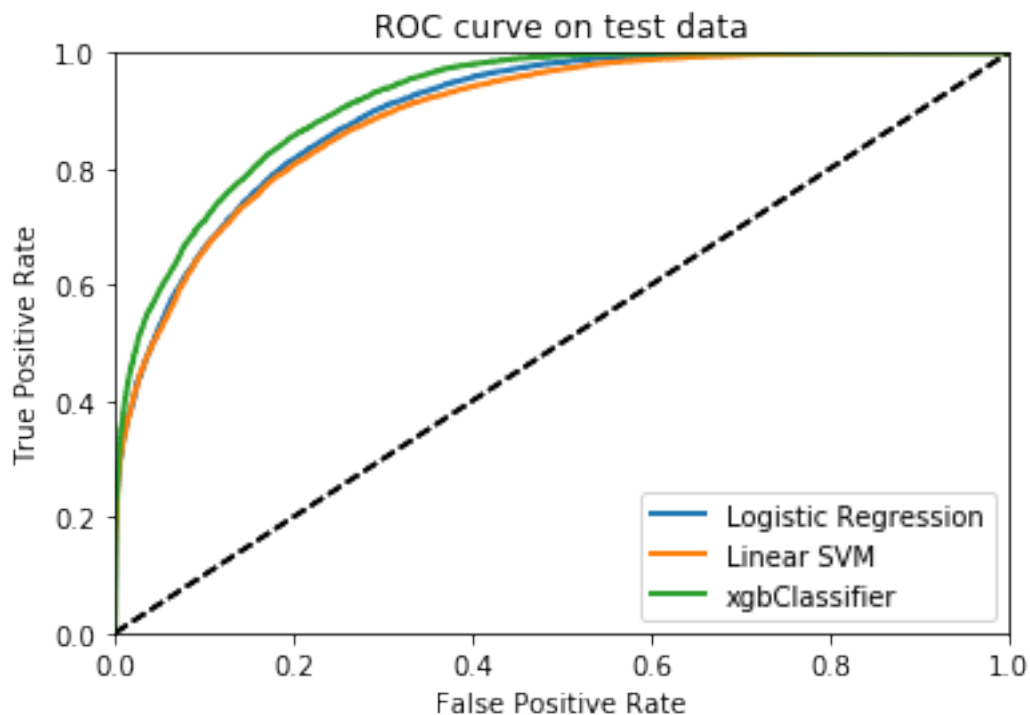
```

AUC score : 91.99

```
In [79]: fpr_BT,tpr_BT,thresholds_BT = roc_curve(Y_test, predict_prob_test[:,1])
```

```
In [83]: def plot_roc_curve(fpr,tpr,label=None):
    plt.plot(fpr,tpr,linewidth=2,label=label)
    plt.plot([0,1],[0,1],"k--")
    plt.xlabel("False Positive Rate")
    plt.ylabel("True Positive Rate")
    plt.axis([0,1,0,1])
    plt.legend()
    plt.title("ROC curve on test data")
```

```
In [84]: plot_roc_curve(fpr_LR,tpr_LR,label='Logistic Regression')
plot_roc_curve(fpr_SVM,tpr_SVM,label='Linear SVM')
plot_roc_curve(fpr_BT,tpr_BT,label='xgbClassifier')
plt.show()
```



```
In [85]: table = PrettyTable(["Model", "Encoding", "Train Log Loss", "Test Log Loss", "Miss cl
```

```
In [86]: table.add_row(['LR', 'TF-IDF', '0.369', '0.385', '18.65', '90.06'])
table.add_row(['Linear SVM', 'TF-IDF', '0.387', '0.395', '18.89', '89.41'])
table.add_row(['xgb Classifier', 'TF-IDF', '0.265', '0.346', '16.29', '92.34'])
```

```
In [87]: print(table)
```



Model	Encoding	Train Log Loss	Test Log Loss	Miss class %	AUC score
LR	TF-IDF	0.369	0.385	18.65	90.06
Linear SVM	TF-IDF	0.387	0.395	18.89	89.41
xgb Classifier	TF-IDF	0.265	0.346	16.29	92.34